

Graph Algorithms with Sparse Linear Algebra

Tim Baer

C3.ai Platform Engineering
University of Illinois at Urbana-Champaign

July 19, 2021

Overview

Introduction

All Pairs Shortest Paths

Sparse Linear Algebra

Minimum Spanning Trees

Expanders

TODO:

- ▶ Introduce notation better/add notation slide?
- ▶ add pauses?
- ▶ Fix footnote numbers

Motivation

Adjacency matrices³ connect graph algorithms with linear algebra.

Designing graph algorithms?

- ▶ Leverage fast matrix multiplication algorithms.
- ▶ Associated matrices have surprising spectral properties.

Designing linear algebra algorithms? (*Not discussed today*)

- ▶ Leverage graph techniques³.


³And Laplacian matrices!

³See *Nearly Linear Time Algorithms for Preconditioning and Solving Symmetric, Diagonally Dominant Linear Systems* by Spielman and Teng.

Adjacency Matrices

Given a graph $G = (V, E)$ with weights $w : E \rightarrow \mathbb{R}$, construct the **adjacency matrix**¹ $\mathbf{A} \in \mathbb{R}^{n \times n}$ where

$$a_{ij} \leftarrow \begin{cases} w(i,j) & \text{if } (i,j) \in E \\ 0 & \text{otherwise} \end{cases}$$



$$\mathbf{A} = \begin{bmatrix} 0 & 5 & 20 & 15 & 0 \\ 5 & 0 & 0 & 5 & 0 \\ 20 & 0 & 0 & 10 & 0 \\ 15 & 5 & 10 & 0 & 10 \\ 0 & 0 & 0 & 10 & 0 \end{bmatrix}$$

¹If G is undirected, store only upper-triangular part of \mathbf{A} .

Algebraic Structures

Monoids (\oplus) generalize *addition*.

$$(\text{"C3"}) + (\text{"ai"}) \rightarrow (\text{"C3.ai"})$$

Semirings (\oplus, \otimes) generalize *addition and multiplication*.

$$x_1 \cdot y_1 + \cdots + x_n \cdot y_n \rightarrow \min\{(x_1 + y_1), \dots, (x_n + y_n)\}$$

Monoids and **semirings** have several nice properties²:

- ▶ Identity elements
- ▶ Associativity
- ▶ ...

²However, neither structure enforces the existence of inverses.

Matrix Multiplication

Matrix-vector multiplication:

$$y \leftarrow \mathbf{A}x$$

Writing as an elementwise expression:

$$y_i \leftarrow \sum_k a_{ik} x_k$$

Generalizing to (\oplus, \otimes) :

$$y_i \leftarrow \bigoplus_k a_{ik} \otimes x_k$$

Matrix-matrix multiplication:

$$\mathbf{C} \leftarrow \mathbf{A}\mathbf{B}$$

Writing as an elementwise expression:


$$c_{ij} \leftarrow \sum_k a_{ik} b_{kj}$$

Generalizing to (\oplus, \otimes) :

$$c_{ij} \leftarrow \bigoplus_k a_{ik} \otimes b_{kj}$$

All Pairs Shortest Paths (APSP)

For each pair of vertices $i, j \in V$, compute the length of a shortest path from i to j .




Sample application: **driving directions**, where road junctions are vertices and road distances are edge weights.

Bellman-Ford Algorithm

Fix a vertex i . Let $d_{ij}^{(\ell)}$ denote the length of a shortest path from i to another vertex j consisting of at most ℓ edges.

An edge (i, j) is **tense** if $d_{ij}^{(\ell)} > w(i, k) + d_{kj}^{(\ell)}$.



For $\ell \leftarrow 1$ to $n - 1$, **relax** all tense edges:

$$d_{ij}^{(\ell+1)} \leftarrow \min_k \{ w(i, k) + d_{kj}^{(\ell)} \}$$

A Linear Algebraic Algorithm

We store tentative shortest path distances in $\mathbf{D}^{(\ell)} \in \mathbb{R}^{n \times n}$ where

$$\mathbf{D}^{(\ell+1)} \leftarrow \mathbf{D}^{(\ell)} \oplus \mathbf{A}\mathbf{D}^{(\ell)} \quad \text{on } (\min, +)$$

Writing as an elementwise expression:

$$d_{ij}^{(\ell+1)} \leftarrow d_{ij}^{(\ell)} \oplus \left(\bigoplus_k a_{ik} \otimes d_{kj}^{(\ell)} \right)$$

Replacing the generic (\oplus, \otimes) with the **tropical semiring** $(\min, +)$:

$$d_{ij}^{(\ell+1)} \leftarrow \min\{d_{ij}^{(\ell)}, \min_k \{w(i, k) + d_{kj}^{(\ell)}\}\}$$

We may accelerate via matrix squaring $\mathbf{D}^{(\ell)} \leftarrow \mathbf{D}^{(\ell/2)} \oplus \mathbf{D}^{(\ell/2)} \mathbf{D}^{(\ell/2)}$.

Matrix Multiplication is Fast

Matrix multiplication has been studied from a variety of algorithmic lenses:

- ▶ Time complexity³
- ▶ Parallel and distributed
- ▶ With sparse data

³The current best time complexity is $O(n^{2.3728596})$ by Alman and Vassilevska via the laser method.

Strassen's Algorithm

$$\begin{bmatrix} \mathbf{C}_{11} & \mathbf{C}_{12} \\ \mathbf{C}_{21} & \mathbf{C}_{22} \end{bmatrix} \leftarrow \begin{bmatrix} \mathbf{A}_{11} & \mathbf{A}_{12} \\ \mathbf{A}_{21} & \mathbf{A}_{22} \end{bmatrix} \begin{bmatrix} \mathbf{B}_{11} & \mathbf{B}_{12} \\ \mathbf{B}_{21} & \mathbf{B}_{22} \end{bmatrix}$$

$$\mathbf{M}_1 \leftarrow (\mathbf{A}_{11} + \mathbf{A}_{22}) \cdot (\mathbf{B}_{11} + \mathbf{B}_{22})$$

$$\mathbf{M}_2 \leftarrow (\mathbf{A}_{21} + \mathbf{A}_{22}) \cdot \mathbf{B}_{11}$$

$$\mathbf{M}_3 \leftarrow \mathbf{A}_{11} \cdot (\mathbf{B}_{12} - \mathbf{B}_{22})$$

$$\mathbf{M}_4 \leftarrow \mathbf{A}_{22} \cdot (\mathbf{B}_{21} - \mathbf{B}_{11})$$

$$\mathbf{M}_5 \leftarrow (\mathbf{A}_{11} + \mathbf{A}_{12}) \cdot \mathbf{B}_{22}$$

$$\mathbf{M}_6 \leftarrow (\mathbf{A}_{21} - \mathbf{A}_{11}) \cdot (\mathbf{B}_{11} + \mathbf{B}_{12})$$

$$\mathbf{M}_7 \leftarrow (\mathbf{A}_{12} - \mathbf{A}_{22}) \cdot (\mathbf{B}_{21} + \mathbf{B}_{22})$$

$$\mathbf{C}_{11} \leftarrow \mathbf{M}_1 + \mathbf{M}_4 - \mathbf{M}_5 + \mathbf{M}_7$$

$$\mathbf{C}_{12} \leftarrow \mathbf{M}_3 + \mathbf{M}_5$$

$$\mathbf{C}_{21} \leftarrow \mathbf{M}_2 + \mathbf{M}_4$$

$$\mathbf{C}_{22} \leftarrow \mathbf{M}_1 - \mathbf{M}_2 + \mathbf{M}_3 + \mathbf{M}_6$$

The time complexity recurrence is given by

$$T(n) = 7T(n/2) + O(n^2) = O(n^{\log_2 7}) \approx O(n^{2.807})$$

Requires the existence of an *additive inverse*.

Many Real-World Graphs are Large

Real-world graphs are **too large** to process on a single machine.

Graph	n	m
2002 crawl of the .uk domain	18.5M	261.8M
Friendster	65.6M	1.8B
Full USA road network	23.9M	28.9M
Orkut social network	3.1M	117M

Table: Publicly available large graphs

We may **distribute** computations across multiple machines³.

³Using a cluster of commodity machines is much cheaper than building a large, specialized machine.

Parallel Matrix Multiplication


TODO

3D Matrix Multiplication was poorly introduced... introduce data layout in another slide (alex what does process owns ... mean?) or introduce parallel/distribute matrix algorithms? Would be nice to show trivial matrix-matrix distribution and why it's bad for communication

3D Matrix Multiplication

TODO, also might create new ideas for multilinear distribution
Organize processes into a 3D cube.

```
for i in range(n):  
    for j in range(n):  
        for k in range(n):  
            c[i,j] += a[i,k] * b[k,j]
```



Process (r, s, t) owns $\mathbf{A}^{(r,t)}$ and $\mathbf{B}^{(t,s)}$ and contributes $\mathbf{C}^{(r,s)}$
where $\mathbf{C}^{(r,s)} \leftarrow \mathbf{A}^{(r,t)} \mathbf{B}^{(t,s)}$.

Figure from CS 484: Parallel Programming.

See *Communication-Optimal Parallel 2.5D Matrix Multiplication and LU Factorization Algorithms* by Solomonik and Demmel (2011).

Many Real-World Graphs are Very Sparse

Many real-world graphs have **relatively few edges**.

Graph	n	m	sp
2002 crawl of the .uk domain	18.5M	261.8M	1.5e-6
Friendster	65.6M	1.8B	8.4e-9
Full USA road network	23.9M	28.9M	1.0e-7
Orkut social network	3.1M	117M	2.4e-5

Table: Publicly available large graphs

However, storing an adjacency matrix requires $O(n^2)$ **entries**.

Compressed Sparse Matrix Formats

We may store only the *nonzero* entries for **sparse** matrices.

Compressed sparse row⁴ (CSR) stores 3 lists:

- ▶ Values (length nnz)
- ▶ Column indices (length nnz): column of value
- ▶ Row offsets (length $n + 1$): number of values above this row

$$\begin{bmatrix} 0 & 1 & 0 & 2 \\ 0 & 0 & 3 & 0 \\ 4 & 0 & 0 & 0 \\ 5 & 6 & 0 & 0 \end{bmatrix}$$

Values: 1, 2, 3, 4, 5, 6

Column indices: 1, 3, 2, 0, 0, 1

Row offset: 0, 2, 3, 4, 6

⁴See also coordinate list and compressed sparse column (CSC) formats.

Compressed Sparse Matrix Algorithms

Sparse matrix-vector multiplication (SpMV):

$$y_i \leftarrow \sum_k a_{ik} x_k$$

```
for i in range(n):  
    for t in range(row_off[i], row_off[i+1]):  
        y[i] += val[t] * x[col_idx[t]]
```

Sparse matrix-matrix multiplication (SpMSpM):

- Communication-optimal algorithm is equivalent to partitioning a hypergraph⁵.

⁵*Hypergraph Partitioning for Sparse Matrix-Matrix Multiplication* by Ballard et al (2016).

Cyclops Tensor Framework

Matrix multiplications can be generalized to **tensors**⁶:

$$u_{i_1 r}^{(1)} \leftarrow \sum_{i_2 \dots i_d} t_{i_1 \dots i_d} u_{i_2 r}^{(2)} \cdots u_{i_d r}^{(d)}$$

The **Cyclops Tensor Framework**⁷ automatically parallelizes custom operations on tensors:

- ▶ Distributed memory
- ▶ Sparse tensors and operations
- ▶ Arbitrary semirings

Tensor algebra is a *framework* for expressing graph algorithms.

⁶For example, MTTKRP is the main kernel for tensor decomposition.

⁷A *Massively Parallel Tensor Contraction Framework for Coupled-Cluster Computations* by Solomonik et. al (2014).

Hybrid Programming (OpenMP + MPI)

The Cyclops Tensor Framework uses **hybrid programming** to leverage parallelism at the *thread* and *process* level.

OpenMP

- ▶ Shared memory
- ▶ Fork, exec, wait model

```
#pragma omp parallel for  
for (i = 0; i < n; i++) {  
    // do something  
}
```

MPI⁸


- ▶ Distributed memory
- ▶ Point-to-point and collectives
- ▶ Standard for HPC

```
// first node  
MPI_Send(&send_data, ...)  
  
// second node  
MPI_Recv(&recv_data, ...)
```

⁸In distributed deep learning, the *AllReduce* model often communicates via NCCL, which is modeled after MPI.

Minimum Spanning Tree (MST)

Find a spanning tree whose total weight is as small as possible.




Sample application: a phone company laying cable in a neighborhood, where houses are vertices and cable lengths are edge weights.

Minimum Spanning Tree Algorithms

Most⁹ algorithms rely only on the **cut property** of MSTs.

Given a cut, the edge with the smallest weight that crosses the cut belongs to the mst.



⁹Karger et. al's $O(m)$ expected time algorithm also uses the **cycle property**.


Awerbuch-Shiloach Parallel MST Algorithm

Maintain a forest of **stars**¹⁰, where each vertex has a parent.

At the first iteration, the forest consists of n vertices, each with a self-loop.


Repeat the following steps until convergence of the forest:

- ▶ Hooking: grows the MST by joining two stars together.
- ▶ Shortcutting: compresses trees into stars.



¹⁰A **star** is a directed tree of height at most 1 and is intuitively a part of the minimum spanning tree.

Forest of Stars




The **parent vector** $p \in \mathbb{R}^n$ stores the parent of each vertex.

The **parent matrix** $P \in \mathbb{R}^{n \times n}$ is defined

$$p_{ij} \leftarrow \begin{cases} 1 & \text{if } p_i = j \\ 0 & \text{otherwise} \end{cases}$$

Hooking

Each star attempts¹¹ to hook with its *smallest cut edge*.




Hooking: joins two stars together.

$$p_i \leftarrow p \left[\underset{\substack{e \in \text{cut}(S_i, V \setminus S_i) \\ \text{edges that cross cut}}}{\text{argmin}} \quad w(e) \right]$$

¹¹Need to break ties to prevent cycles.

Shortcut

Hooking may generate trees of arbitrary height.



Shortcutting¹²: reduces the height of each tree in the forest by a factor of nearly two. $p_i \leftarrow p_{p_i}$

¹²Until completion so that all trees become stars.

Hooking as a Multilinear Operation

f computes for each pair of nodes (i, j) , whether they belong to the same star.

$$f(p_i, a_{ij}, p_j) \leftarrow \begin{cases} a_{ij} : p_i \neq p_j \\ \infty : \text{otherwise} \end{cases}$$

Accumulating $f(p_i, a_{ij}, p_j)$ over j computes for each node i , its smallest cut edge.

$$q_i \leftarrow \bigoplus_j f(p_i, a_{ij}, p_j)$$

For each star root i , we **contract** its children's smallest cut edges into itself.


$$r_{p_j} \xleftarrow{\text{MinWeight}} q_j$$

Data Distribution

How can we leverage the structure of $f(p_i, a_{ij}, p_j)$?

$$w_i \leftarrow \bigoplus_j f(x_i, a_{ij}, y_j)$$

Process (r, s) owns $\mathbf{A}^{(r,s)}$ and needs $x^{(r)}$ and $y^{(s)}$. Process (r, s) contributes $w^{(r)}$ where $w^{(r)} \leftarrow f(x^{(r)}, \mathbf{A}^{(r,s)}, y^{(s)})$.




Contraction

To reduce the size of the graph in the next iteration, we may merge vertices via **contraction**.

$\mathbf{C} \leftarrow \mathbf{P}^T \mathbf{A} \mathbf{P}$ on the (\min, \cdot) semiring performs a contraction of \mathbf{A} .

$$\begin{aligned} c_{ij} &\leftarrow \sum_l p_{li}^T \left(\sum_k a_{lk} p_{kj} \right) \\ &= \sum_l p_{li} \left(\sum_k a_{lk} p_{kj} \right) \\ &= \sum_{l,k} p_{li} a_{lk} p_{kj} \end{aligned}$$

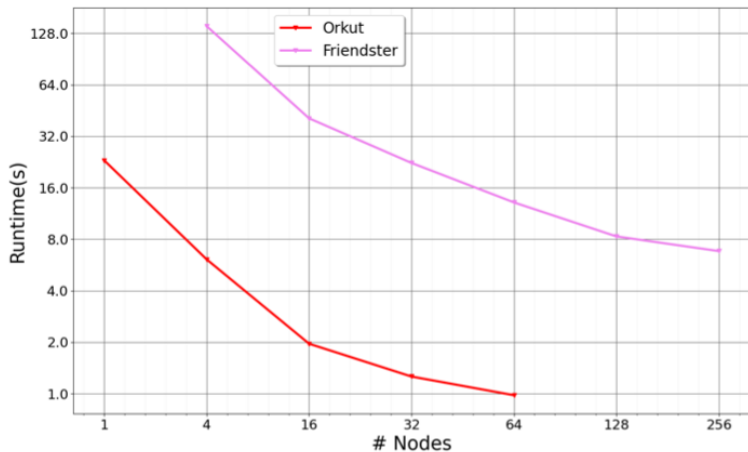


$w(l, k)$ is projected onto $w(p_l, p_k)$ if:

- ▶ The edge (l, k) exists
- ▶ i is l 's parent ($p_{li} = 1$)
- ▶ j is k 's parent ($p_{kj} = 1$)

MST Strong Scaling Data


Strong scaling: fix the *total* problem size and vary the # of nodes.



(a) Social network graphs

MST Weak Scaling Data

Weak scaling: fix the problem size on *each* node and vary the # of nodes & total problem size accordingly.




(a) Uniform random graph

Figure 2: Edge weak scaling of uniform random graphs.
Constant $n^2/p = n_0^2$ and edge percentage $f = 100 * m/n^2$

Markov Chains

Markov chain: Given a set of states S , if the current state is i then go to state j with probability p_{ij} .



$$P = \begin{bmatrix} 0.6 & 0.3 & 0.1 \\ 0.2 & 0.3 & 0.5 \\ 0.4 & 0.1 & 0.5 \end{bmatrix}$$

Random Walks

We prove expectations of the below via analyzing **random walks**:

- ▶ **Hitting time**: number of steps to reach state j from state i ?
- ▶ **Cover time**: number of steps to reach all states from state i ?
- ▶ **Stationary distribution**: what is the probability of ending in state i in a walk with t steps as $t \rightarrow \infty$?
- ▶ **Mixing time**: number of steps to approach the stationary distribution?

A vector $\pi = (\pi)_{i \in S}$ is the **stationary distribution** if

$$\pi = \pi \mathbf{P}$$

The stationary distribution is an eigenvector of \mathbf{P} with eigenvalue 1.

Mixing Time

From a starting distribution $\pi^{(0)}$, take a random walk with t steps.

$$\begin{aligned}\pi^{(t)} &\leftarrow \pi^{(t-1)} \mathbf{P} \\ &\vdots \\ &= \pi^{(0)} \mathbf{P}^t\end{aligned}$$

Mixing time: what is the expected number of steps to approach the stationary distribution?

The above algorithm is the same as **power iteration**¹³:

- ▶ Converges to the dominant eigenvector of \mathbf{P} (namely, π)
- ▶ Convergence ratio is given by $|\lambda_1/\lambda_2|$

¹³Instead of normalization, enforce that $\sum_i \pi_i = 1$.

Expanders

A d -regular graph is an α -expander¹⁴ if $|\lambda_2| \leq \dots \leq |\lambda_n| \leq \alpha$.

Rapid mixing lemma: For an α -expander,

$$(\pi^{(t)} - \pi) \leq \alpha^t$$

Do good expanders ($\alpha = O(\frac{1}{\sqrt{d}})$) exist?

- ▶ Random graph (probabilistic method)
- ▶ Graph products (zig-zag product)
- ▶ Ramanujan graphs ($\alpha = \frac{1}{\sqrt{d}}$)

¹⁴ $\alpha < 1$. Combinatorial definition: A d -regular graph is an α -expander if

$$\frac{\partial S}{|S|} \geq \Omega(d(1 - \alpha)) \quad \forall S \subseteq V \quad (|S| \leq n/2)$$

Applications of Expanders

Expanders are a very powerful tool in modern TCS:

- ▶ Derandomization (**randomness** is a *resource*)
- ▶ Error correcting codes
- ▶ Sorting networks
- ▶ ...





Figure from *Cutoff on All Ramanujan Graphs* by Lubetzky and Peres (2015).

Monte Carlo Algorithms

A **Monte Carlo algorithm** Π for a decision problem:

- ▶ Correctly outputs YES with probability $\geq 1/2$
- ▶ Correctly outputs NO with probability 1

Is the polynomial¹⁵ $P(x_1, \dots, x_k)$ over \mathbb{Z}_p **zero everywhere**?



P has at most degree d roots:

- ▶ Pick a random vector $x_1, \dots, x_k \in \mathbb{Z}_p$.
- ▶ $\Pr[P(x_1, \dots, x_k) = 0] \leq \frac{d}{p}$.
- ▶ If YES, repeat?
- ▶ If NO, output NO.

¹⁵For example, $((x^2 - 1)(y^2 - 4) + x^2 + y^2 - 5)/(x^2 + y^2 + 1)^2$.

Figure from *Visualizing Functions of Several Variables and Surfaces* by

Probability Amplification

Consider a Monte Carlo algorithm Π :

- ▶ Error probability $\leq 1/10$.
- ▶ m random bits.

How can we *lower* the error probability?

Repeat t times \rightarrow error probability $\leq (\frac{1}{10})^t$ and tm random bits.

Can we achieve a similar error probability with fewer random bits?

- ▶ Expanders!


Derandomization

Construct a d -regular α -expander G with $n = 2^m$ vertices¹⁶.

Fix a subset $B \subseteq V$ with $|B| = \beta n$.

Take a random walk v_1, \dots, v_t from a random starting vertex.

$$\Pr[v_1 \in B \wedge \dots \wedge v_t \in B] \leq (\alpha + \beta)^t$$



¹⁶Efficient representations have been designed to construct G “on the fly”.

Derandomization (cont.)

If G is a good expander, we are *unlikely* to get stuck in B ¹⁷.

Probability amplification:

- ▶ Each vertex corresponds to a m bit string.
- ▶ When evaluating Π , we randomly pick a m bit string.

B is the set of m bit strings that output a wrong answer to Π .

$$|B| \leq \frac{1}{10}|n| \implies \beta \leq \frac{1}{10}$$

- ▶ Error probability $\leq (\alpha + \frac{1}{10})^t$.
- ▶ $m + t \log d$ random bits.

¹⁷Can also be viewed from the combinatorial definition of expanders.

Summary

Introduction

All Pairs Shortest Paths

Sparse Linear Algebra

Minimum Spanning Trees

Expanders

References

- ▶ *Graph Algorithms in the Language of Linear Algebra* by Kepner and Gilbert (2011)
- ▶ CS 554: Parallel Numerical Algorithms by Solomonik
- ▶ *New Connectivity and MSF Algorithms for Shuffle-Exchange Network and PRAM* by Awerbuch and Shiloach (1987)
- ▶ CS 574: Randomized Algorithms by Chan

Fin

Thanks. Questions?

Extra Slides

The following are cool related topics that I'd like to talk about at some point:

- ▶ Ball: given a vertex, compute its b closest neighbors.
- ▶ Probabilistic method: prove existence then derandomize for a construction.
- ▶ Dimensionality reduction: does there exist an embedding from ℓ_1 to Hamming space? Take a random projection and use the probabilistic method.
- ▶ Convolution as matrix multiplication: use Toeplitz matrices.
- ▶ Tensor decomposition: generalizations of SVD.
- ▶ Tensor completion: Netflix Prize.
- ▶ Hardware accelerators for matrix multiplications.