

Low Hop Emulators and Uncapacitated Min-Cost Flow¹

Caleb Ju and Tim Baer

Parallel Algorithms Reading Group
University of Illinois at Urbana-Champaign

October 15, 2020

¹Based on results from *Parallel Approximate Undirected Shortest Paths Via Low Hop Emulators* by Andoni, Stein, and Zhong (published in STOC 2020).

- 1 Low Hop Emulators and Applications
- 2 Constructing a Low Hop Emulator
- 3 Uncapacitated Min-Cost Flow in Sherman's Framework
- 4 Constructing a Good Preconditioner

Recent Results

Theorem (Andoni et. al SSSP)

$\text{polylog}(n)$ -approximate single source shortest path with $\text{polylog}(n)$ depth and $m \cdot \text{polylog}(n)$ work via *low hop emulators*.

Theorem (Andoni et. al (s, t) -SP)

$(1 + \epsilon)$ -approximate $(s - t)$ -shortest path with $\text{polylog}(n)$ depth and $m \cdot \text{polylog}(n)$ work via reduction to *uncapacitated min-cost flow*.

²*Faster Parallel Algorithm for Approximate Shortest Path* by Li (published in STOC 2020).

Recent Results

Theorem (Andoni et. al SSSP)

$\text{polylog}(n)$ -approximate single source shortest path with $\text{polylog}(n)$ depth and $m \cdot \text{polylog}(n)$ work via *low hop emulators*.

Theorem (Andoni et. al (s, t) -SP)

$(1 + \epsilon)$ -approximate $(s - t)$ -shortest path with $\text{polylog}(n)$ depth and $m \cdot \text{polylog}(n)$ work via reduction to *uncapacitated min-cost flow*.

Theorem (Li SSSP²)

$(1 + \epsilon)$ -approximate single source shortest path with $\text{polylog}(n)$ depth and $m \cdot \text{polylog}(n)$ work via reduction to *uncapacitated min-cost flow*.

²Faster Parallel Algorithm for Approximate Shortest Path by Li (published in STOC 2020).

What is a Low Hop Emulator?

What is a **low hop emulator**?

Given a graph $G = (V, E)$, a **low hop emulator** is a weighted graph $H = (V, F)$ where any shortest $(s - t)$ -path uses $\mathcal{O}(\log \log n)$ edge traversals and $|F| = \mathcal{O}(m \cdot \text{poly}(\log n))$.

³Uses hierarchy of discretizations to solve a linear system.

⁴ $d(u, v) \leq \tilde{d}(u, v) \leq f(n)d(u, v)$, where $f(n)$ is the approximation factor, and d and \tilde{d} are the exact and approximate distances, respectively.

What is a Low Hop Emulator?

What is a **low hop emulator**?

Given a graph $G = (V, E)$, a **low hop emulator** is a weighted graph $H = (V, F)$ where any shortest $(s - t)$ -path uses $\mathcal{O}(\log \log n)$ edge traversals and $|F| = \mathcal{O}(m \cdot \text{poly}(\log n))$.

- Is multigrid inspired³
- Provable approximation/distortion factors⁴
- Can be constructed in parallel

³Uses hierarchy of discretizations to solve a linear system.

⁴ $d(u, v) \leq \tilde{d}(u, v) \leq f(n)d(u, v)$, where $f(n)$ is the approximation factor, and d and \tilde{d} are the exact and approximate distances, respectively.

- ① Uncapacitated min-cost flow: find cheapest way to *route* flow from supply vertices to demand vertices
- ② Bourgain's embedding: embed *any* metric space into ℓ_p with distortion $\mathcal{O}(\log n)$
- ③ Low diameter decomposition: decompose a graph into subsets such that *far* vertices are *unlikely* to belong to the same subset

Constructing the Low Hop Emulator

Given a graph $G = (V, E)$, a **low hop emulator** is a weighted graph $H = (V, F)$ where any shortest $(s - t)$ -path uses $\mathcal{O}(\log \log n)$ edge traversals and $|F| = \mathcal{O}(m \cdot \text{poly}(\log n))$.

How can we build a low hop emulator?

Constructing the Low Hop Emulator

Given a graph $G = (V, E)$, a **low hop emulator** is a weighted graph $H = (V, F)$ where any shortest $(s - t)$ -path uses $\mathcal{O}(\log \log n)$ edge traversals and $|F| = \mathcal{O}(m \cdot \text{poly}(\log n))$.

How can we build a low hop emulator?

- 1 Build a sparser graph using a **subemulator**
- 2 Recursively apply subemulator $\mathcal{O}(\log \log n)$ times

Constructing the Low Hop Emulator

Given a graph $G = (V, E)$, a **low hop emulator** is a weighted graph $H = (V, F)$ where any shortest $(s - t)$ -path uses $\mathcal{O}(\log \log n)$ edge traversals and $|F| = \mathcal{O}(m \cdot \text{poly}(\log n))$.

How can we build a low hop emulator?

- 1 Build a sparser graph using a **subemulator**
- 2 Recursively apply subemulator $\mathcal{O}(\log \log n)$ times

A **subemulator** is a graph $H = (S, F')$ where $S \subset V$ and F' is a weighted edge set that approximates distances well.

Constructing a Subemulator

A **subemulator** is a graph $H = (S, F')$ where $S \subset V$ and F' is a weighted edge set that approximates distances well.

- ① Initially construct S by sampling vertices
- ② Add more vertices to ensure every vertex is close to a vertex in S
- ③ Add weighted edges so that local distances are well-approximated

Constructing a Subemulator

A **subemulator** is a graph $H = (S, F')$ where $S \subset V$ and F' is a weighted edge set that approximates distances well.

- 1 Initially construct S by sampling vertices
- 2 Add more vertices to ensure every vertex is close to a vertex in S
- 3 Add weighted edges so that local distances are well-approximated

The **ball**, $B_{G,b}(v)$: closest b vertices (w.r.t graph distance) to v .

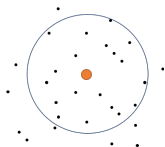


Figure: Ball

Selecting Vertices

- 1 Initially construct S by sampling vertices
- 2 Add vertices to S so that every $v \in V$ is close to a vertex in S
- 3 Add weighted edges so that local distances are well-approximated

Given ball size b (typically $b = \mathcal{O}(\log n)$).

- 1 For every vertex $v \in V$, construct its ball $B(v)$
- 2 Construct S by sampling every vertex with probability $p = \min(50 \frac{\log n}{b}, \frac{1}{2})$
- 3 For any $v \in V \setminus S$ whose ball does not contain any vertex in S , then add v to S
- 4 Store the $\text{leader}(v) \leftarrow$ closest vertex $u \in S$ to $v \in V$

Output: A sparse vertex set S and leader mapping $q : V \rightarrow S$

Selecting vertices

Can we only compute the balls for the **sampled** vertices?

For any $v \in V \setminus S$ whose ball does not contain any vertex in S , then add v to S




Figure: Line 4 of selecting vertices

For any $v \in V \setminus S$ that is not contained in the ball of any vertex in S , then add v to S




Figure: Only compute balls for sampled vertices

Adding Edges

- ➊ Initially construct S by sampling vertices
- ➋ Add more vertices to ensure vertex in V is close to a vertex in S
- ➌ Add weighted edges so that local distances are well-approximated

Adding Edges

- 1 Initially construct S by sampling vertices
- 2 Add more vertices to ensure vertex in V is close to a vertex in S
- 3 Add weighted edges so that local distances are well-approximated

Denote $q(v) = \text{leader}(v) \in S$




Figure: For every edge from G , add edge between the leaders

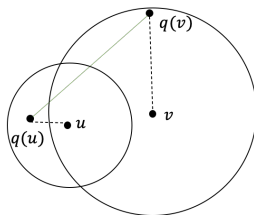


Figure: For every pair of vertices that are “close”, add an edge between their leaders

Adding Edges

Denote $q(v) = \text{leader}(v) \in S$.

- 1 For every $(u, v) \in E$, add edge $(q(u), q(v))$
- 2 For every $v \in V, u \in B(v)$, add edge $(q(u), q(v))$

Set $w(e) = \min \begin{cases} w(e), (\text{initialize to } \infty) \\ d_G(q(u), u) + d_G(u, v) + d_G(v, q(v)) \end{cases}$.

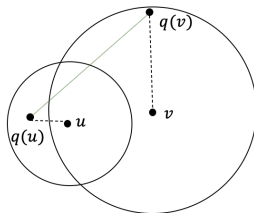
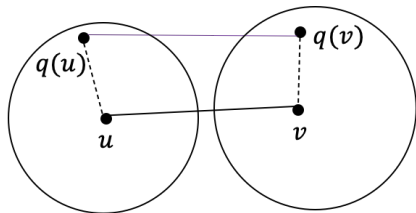


Figure:

Properties of Subemulators

Let $H = (S, F')$ be a subemulator of G .

- $\mathbb{E}[|S|] \leq \frac{3}{4}n$ and $|F'| \leq m + nb$
- For any $u, v \in S$, their distance in H is distorted by a constant factor⁵
- For any $u, v \in V$, the distance between their leaders in H is distorted by a constant factor⁶

A **subemulator** is a graph $H = (S, F')$ where $S \subset V$ and F' is a weighted edge set that approximates distances well.

⁵For any $u, v \in S$, $d_G(u, v) \leq d_H(u, v) \leq 8 \cdot d_G(u, v)$.

⁶For any $u, v \in V$, $d_H(q(u), q(v)) \leq d_G(u, q(u)) + 22 \cdot d_G(u, v) + d_G(v, q(v))$.

Constructing the Low Hop Emulator

Given a graph $G = (V, E)$, a **low hop emulator** is a weighted graph $H = (V, F)$ where any shortest $(s - t)$ -path uses $\mathcal{O}(\log \log n)$ edge traversals and $|F| = \mathcal{O}(m \cdot \text{poly}(\log n))$.

How can we build a low hop emulator?

- 1 Define a restriction operator via subemulator
- 2 Recursively apply subemulator $\mathcal{O}(\log \log n)$ times

Combining Nested Subemulators as H (Distance Oracle)

- 1 Form $t = \mathcal{O}(\log \log n)$ recursive subemulators
- 2 Add an edge from a vertex to its leaders in the next level
- 3 Keep edges within each subemulator
- 4 Add edges between “close” vertices in the same level

Scale: $w_H(u, v) = 27^{t - \max(|\text{lvl}(u), \text{lvl}(v)|)} \cdot d_{H_i}(u, v).$

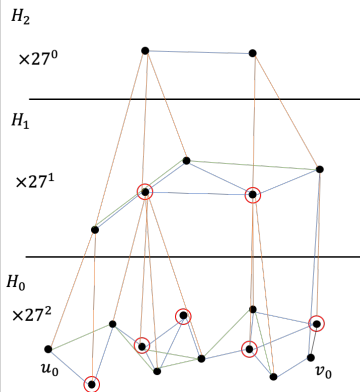


Figure: Scaling edges

Properties of a Distance Oracle

Lemma

Let H be distance oracle.

- $|E(H)| = \mathcal{O}(m \cdot \text{poly}(\log n))$
- *For every $u, v \in V$ and corresponding $u^{(0)}, v^{(0)} \in H_0$ (in the bottom level of H), then $d_H(u^{(0)}, v^{(0)})$ is distorted by at most $\text{polylog}(n)$ ⁷*

⁷Precisely, $d_G(u, v) \leq d_H(u^{(0)}, v^{(0)}) \leq 26^{\mathcal{O}(\log \log n)} \cdot d_G(u, v)$

Properties of a Distance Oracle

Lemma

Let H be distance oracle.

- $|E(H)| = \mathcal{O}(m \cdot \text{poly}(\log n))$
- *For every $u, v \in V$ and corresponding $u^{(0)}, v^{(0)} \in H_0$ (in the bottom level of H), then $d_H(u^{(0)}, v^{(0)})$ is distorted by at most $\text{polylog}(n)^7$*

Given a graph $G = (V, E)$, a **low hop emulator** is a weighted graph $H = (V, F)$ where any shortest $(s - t)$ -path uses $\mathcal{O}(\log \log n)$ edge traversals and $|F| = \mathcal{O}(m \cdot \text{poly}(\log n))$.

⁷Precisely, $d_G(u, v) \leq d_H(u^{(0)}, v^{(0)}) \leq 26^{\mathcal{O}(\log \log n)} \cdot d_G(u, v)$

Creating the Ball Centered at v (in parallel)

$B_{G,b}(v) = \{\text{closest } b \text{ vertices to } v\}$.

Goal: Build all n balls in $\log n$ depth and nearly linear work.

Creating the Ball Centered at v (in parallel)

$B_{G,b}(v) = \{\text{closest } b \text{ vertices to } v\}.$

Goal: Build all n balls in $\log n$ depth and nearly linear work.

Idea 1: Compute all pairs shortest path (APSP) using $\log(n)$ *path doublings*. Then save b closest vertices.

Cost: $\log(n)$ iterations with $\mathcal{O}(n^3 \cdot \log(n))$ work.

Creating the Ball Centered at v (in parallel)

$B_{G,b}(v) = \{\text{closest } b \text{ vertices to } v\}$.

Goal: Build all n balls in $\log n$ depth and nearly linear work.

Idea 1: Compute all pairs shortest path (APSP) using $\log(n)$ *path doublings*. Then save b closest vertices.

Cost: $\log(n)$ iterations with $\mathcal{O}(n^3 \cdot \log(n))$ work.

Idea 2: Compute *partial* APSP using $\log(b)$ *truncated* path doublings. After each path doubling, keep b closest vertices.

Cost: $\log(b)$ iterations with $\mathcal{O}(n \cdot b^2 \cdot \log(n))$ work. Set $b = \mathcal{O}(\log n)$.

Creating the Ball Centered at v (in parallel)

Let $\mathbf{B}^{(i)} \in \mathbb{R}^{n \times n}$ on the $(\min, +)$ semiring contain the tentative distances to the b closest vertices to vertex $u, \forall u \in V$.

$$\mathbf{B}^{(i+1)} = F[\mathbf{B}^{(i)} \mathbf{B}^{(i)}]$$

where F extracts the b closest neighbors to $u, \forall u \in V$.

Creating the Ball Centered at v (in parallel)

Let $\mathbf{B}^{(i)} \in \mathbb{R}^{n \times n}$ on the $(\min, +)$ semiring contain the tentative distances to the b closest vertices to vertex $u, \forall u \in V$.

$$\mathbf{B}^{(i+1)} = F[\mathbf{B}^{(i)} \mathbf{B}^{(i)}]$$

where F extracts the b closest neighbors to $u, \forall u \in V$.

- ① Requires $\mathcal{O}(\log b)$ iterations
- ② Parallelizable using existing matrix-matrix product techniques
- ③ Can be sparsified by specifying a row filter

Creating the Ball Centered at v (in parallel)

Alternatively, let $B_u^{(i)}$ be a sorted b -vector containing tentative distances to the b closest vertices to vertex u .

$$B_u^{(i+1)} = B_u^{(i)} \oplus \left(\bigoplus_{(u,v) \in E} \{w(u,v) + B_v^{(i)}\} \right)$$

where the reduction operator $x \oplus y$ merges x and y and returns the first b distances.

Creating the Ball Centered at v (in parallel)

Alternatively, let $B_u^{(i)}$ be a sorted b -vector containing tentative distances to the b closest vertices to vertex u .

$$B_u^{(i+1)} = B_u^{(i)} \oplus \left(\bigoplus_{(u,v) \in E} \{w(u,v) + B_v^{(i)}\} \right)$$

where the reduction operator $x \oplus y$ merges x and y and returns the first b distances.

- 1 Requires $\mathcal{O}(b)$ iterations
- 2 Can be embedded into a generalized **matrix-vector product** with the adjacency matrix and parallelized with existing techniques
- 3 Can be sparsified by specifying a row filter

Tuning Parameters

- ① approximation constant ϵ
- ② size of the ball b
- ③ number of subemulators
- ④ sampling probability of vertices
- ⑤ scaling factor for combining nested subemulators

Uncapacitated Min–Cost Flow (Transshipment)

Let $\mathbf{W} \in \mathbb{R}^{m \times m}$ be a diagonal matrix of weights. Let $\mathbf{A} \in \mathbb{R}^{n \times m}$ be the incidence matrix,

$$\mathbf{A}_{iu} = \begin{cases} 1 & : \exists \text{ edge } u = (i, j) \\ -1 & : \exists \text{ edge } u = (j, i) \\ 0 & : \text{otherwise} \end{cases}.$$

Find a vector $f \in \mathbb{R}^m$ such that

$$\begin{aligned} \min_{f \in \mathbb{R}^m} & \quad \|\mathbf{W}f\|_1 \\ \text{s.t.} & \quad \mathbf{A}f = b, \end{aligned}$$

where $b \in \mathbb{R}^n$ is the **demand vector**, where we require $\sum_i b_i = 0$.

If $b(s) = 1$, $b(t) = -1$, then solves (s, t) –shortest path length.

Uncapacitated Min-Cost Flow (Transshipment)

An equivalent problem:

Let $x = \mathbf{W}f$. Find the optimal x^* such that

$$\begin{aligned} x^* = \min_{x \in \mathbb{R}^m} \|x\|_1 \\ \text{s.t. } \mathbf{A}\mathbf{W}^{-1}x = b. \end{aligned}$$

Uncapacitated Min–Cost Flow (Transshipment)

An equivalent problem:

Let $x = \mathbf{W}f$. Find the optimal x^* such that

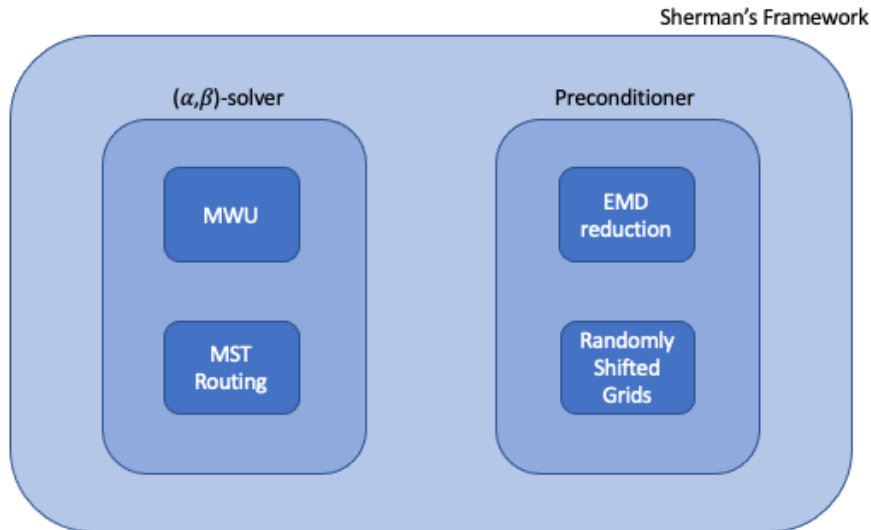
$$\begin{aligned} x^* &= \min_{x \in \mathbb{R}^m} \|x\|_1 \\ \text{s.t. } \mathbf{A}\mathbf{W}^{-1}x &= b. \end{aligned}$$

Lemma: There exists $(1 + \varepsilon)$ -approximation algorithm to the optimization problem above in that runs in polylog depth if there exists a matrix \mathbf{P} such that

$$\|x^*\|_1 \leq \|\mathbf{P}b\|_1 \leq \mathcal{O}(\text{poly log } n) \cdot \|x^*\|_1.$$

Sherman's Framework

We use Sherman's framework to solve uncapacitated min-cost flow.



Sherman's Framework⁸

Let \mathcal{X}, \mathcal{Y} be finite dimensional vector spaces, where \mathcal{X} is also a Banach space, and let $\mathbf{A} \in \text{Lin}(\mathcal{X}, \mathcal{Y})$ be fixed. Consider the problem:

$$\begin{aligned} \min_{x \in \mathcal{X}} \|x\|_{\mathcal{X}} \\ \text{s.t. } \mathbf{A}x = b, \end{aligned}$$

⁸Based on results from *Generalized Preconditioning and Network Flow Problems* by Sherman (published in SODA 2017).

Sherman's Framework⁸

Let \mathcal{X}, \mathcal{Y} be finite dimensional vector spaces, where \mathcal{X} is also a Banach space, and let $\mathbf{A} \in \text{Lin}(\mathcal{X}, \mathcal{Y})$ be fixed. Consider the problem:

$$\begin{aligned} \min_{x \in \mathcal{X}} \|x\|_{\mathcal{X}} \\ \text{s.t. } \mathbf{A}x = b, \end{aligned}$$

x is an (α, β) -solution to the above problem if

① $\frac{\|x\|}{\|x_{opt}\|} \leq \alpha$, and

② $\frac{\|\mathbf{A}x - b\|}{\|\mathbf{A}\| \|x_{opt}\|} \leq \beta$

⁸Based on results from *Generalized Preconditioning and Network Flow Problems* by Sherman (published in SODA 2017).

Sherman's Framework

x is an (α, β) -solution to the above problem if

① $\frac{\|x\|}{\|x_{opt}\|} \leq \alpha$, and

② $\frac{\|Ax - b\|}{\|A\| \|x_{opt}\|} \leq \beta$

The **non-linear condition number** of $A : \mathcal{X} \rightarrow \mathcal{Y}$ is

$$\kappa_{\mathcal{X} \rightarrow \mathcal{Y}}(A) = \min \left\{ \frac{\|A\|_{\mathcal{X} \rightarrow \mathcal{Y}} \|x\|_{\mathcal{X}}}{\|Ax\|_{\mathcal{Y}}} : Ax \neq 0 \right\}$$

Sherman's Framework

x is an (α, β) -solution to the above problem if

① $\frac{\|x\|}{\|x_{opt}\|} \leq \alpha$, and

② $\frac{\|Ax - b\|}{\|A\| \|x_{opt}\|} \leq \beta$

The **non-linear condition number** of $A : \mathcal{X} \rightarrow \mathcal{Y}$ is

$$\kappa_{\mathcal{X} \rightarrow \mathcal{Y}}(A) = \min \left\{ \frac{\|A\|_{\mathcal{X} \rightarrow \mathcal{Y}} \|x\|_{\mathcal{X}}}{\|Ax\|_{\mathcal{Y}}} : Ax \neq 0 \right\}$$

Theorem (Composition of solvers)

Let F_i be a $(\alpha_i, \beta_i/\kappa)$ -solver for $A : \mathcal{X} \rightarrow \mathcal{Y}$, where A has non-linear condition number κ . Then, the composition $F_2 \circ F_1$ is an $(\alpha_1 + \alpha_2\beta_1, \beta_1\beta_2/\kappa)$ -solver for the same problem.

Multiplicative Weights Update

We can reduce the problem

$$\begin{aligned} \min_{x \in \mathcal{X}} \|x\|_{\mathcal{X}} \\ \text{s.t. } \mathbf{A}x = b, \end{aligned}$$

to a feasibility problem, which we can approximately solve with the **multiplicative weights update method**.

Multiplicative Weights Update

We can reduce the problem

$$\begin{aligned} \min_{x \in \mathcal{X}} \|x\|_{\mathcal{X}} \\ \text{s.t. } \mathbf{A}x = b, \end{aligned}$$

to a feasibility problem, which we can approximately solve with the **multiplicative weights update method**.

Consider an arbitrary decision and a panel of n experts. Over a series of rounds, the **multiplicative weights update method** rewards experts whose predictions are good and punishes experts whose predictions are poor.

As an iterative method, we can use the **composition of solvers theorem**.

Earth Mover's Distance Problem

Theorem (Bourgain's Embedding)

Given a graph $G = (V, E)$ and distance $d : V \times V \rightarrow \mathbb{R}^+$, there exists a mapping $\phi : V \rightarrow [\Delta]^{\mathcal{O}(\log^2 n)}$ such that

$$d_G(u, v) \leq \|\phi(u) - \phi(v)\|_1 \leq \mathcal{O}(\log n) d_G(u, v),$$

Earth Mover's Distance Problem

Theorem (Bourgain's Embedding)

Given a graph $G = (V, E)$ and distance $d : V \times V \rightarrow \mathbb{R}^+$, there exists a mapping $\phi : V \rightarrow [\Delta]^{\mathcal{O}(\log^2 n)}$ such that

$$d_G(u, v) \leq \|\phi(u) - \phi(v)\|_1 \leq \mathcal{O}(\log n) d_G(u, v),$$

We can reduce approximating the uncapacitated min-cost flow on G to approximating the cost of the **Earth Mover's Distance problem**.

The Earth Mover's Distance (EMD) problem is

$$\begin{aligned} \min_{\pi: V \times V \rightarrow \mathbb{R}_{\geq 0}} \quad & \sum_{(u,v) \in V \times V} \pi(u, v) \cdot \|\phi(u) - \phi(v)\|_1 \\ \text{s.t.} \quad & \forall u \in V, \sum_{v \in V} \pi(u, v) - \sum_{v \in V} \pi(v, u) = b_u. \end{aligned}$$

Randomly Shifted Grids

We use randomly shifted grids to obtain a β -approximation to OPT_{EMD} .

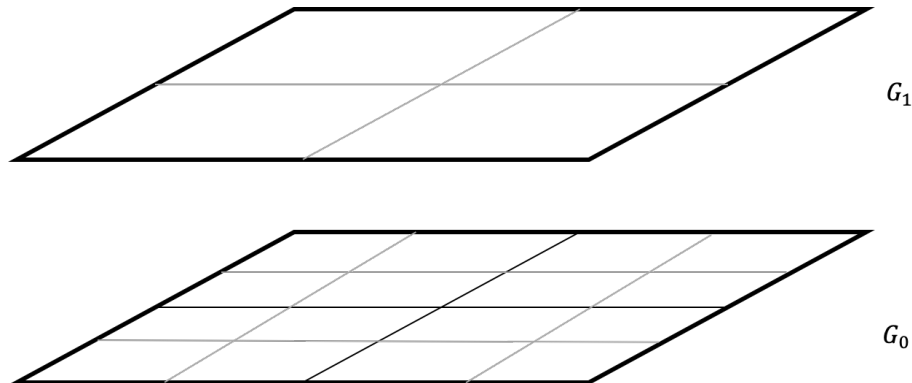


Figure: Grids

Randomly Shifted Grids




Figure: Grids with points

Preconditioner

Construct a sequence of $L = 1 + \log \Delta$ grids G_i and let τ be a random variable over $[\Delta]$.

For each level $i \in \{0, 1, \dots, L-1\}$, each cell $C \in G_i$, and each shift value $\tau \in [2^i]$, we set $h' \in \mathbb{R}^{\sum_{i=0}^{L-1} 2^i |G_i|}$ to

$$h'_{(i,C,\tau)} = d \cdot \sum_{v \in V: \phi(v) + \tau \cdot \mathbf{1}_d \in C} b_v$$

Then,

- ① $OPT_{EMD}(b) \leq \|h'\|_1 \leq 2Ld \cdot OPT_{EMD}(b)$
- ② $\kappa(P'AW^{-1}) \leq 2Ld\alpha$

Preconditioner

Construct a sequence of $L = 1 + \log \Delta$ grids G_i and let τ be a random variable over $[\Delta]$.

For each level $i \in \{0, 1, \dots, L-1\}$, each cell $C \in G_i$, and each shift value $\tau \in [2^i]$, we set $h' \in \mathbb{R}^{\sum_{i=0}^{L-1} 2^i |G_i|}$ to

$$h'_{(i,C,\tau)} = d \cdot \sum_{v \in V: \phi(v) + \tau \cdot 1_d \in C} b_v$$

Then,

- ① $OPT_{EMD}(b) \leq \|h'\|_1 \leq 2Ld \cdot OPT_{EMD}(b)$
- ② $\kappa(P'AW^{-1}) \leq 2Ld\alpha$

Observe that h' can be written as a linear map $h' = Pb$ where

$$P'(i, C, \tau), v = \begin{cases} d & \phi(v) + \tau \cdot 1_d \in C \\ 0 & \text{otherwise} \end{cases}.$$

Preconditioner

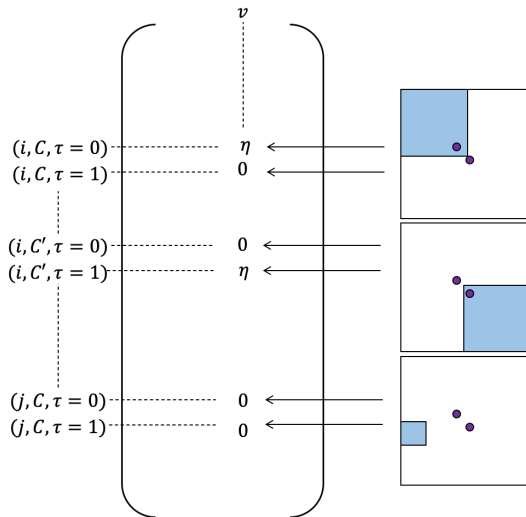
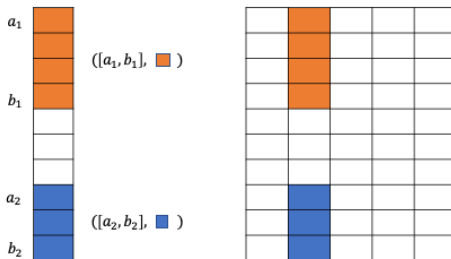


Figure: Preconditioner

Compressed Representation

Given a vector $x \in \mathbb{R}^r$, a **compressed representation** of x is a set of tuples $I = \{([a_1, b_1], c_1), ([a_2, b_2], c_2), \dots, ([a_s, b_s], c_s)\}$ where $c_i \in \mathbb{R}$, $[a_i, b_i] \subseteq [1, r]$ such that

- 1 $\forall i \neq j \in [s], [a_i, b_i] \cap [a_j, c_j] = \emptyset$
- 2 $\forall j \in [a_i, b_i], x_j = c_i$
- 3 $\forall j \in [1, r] \setminus \bigcup_{i \in [s]} [a_i, b_i], x_j = 0$



Implicit Preconditioner

How can we construct our implicit preconditioner?

$$P'(i, C, \tau), v = \begin{cases} d & \phi(v) + \tau \cdot 1_d \in C \\ 0 & \text{otherwise} \end{cases}$$

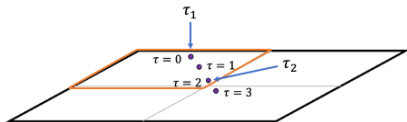


Figure: Grid with shifted points

Fix a vertex $v \in V$, a level $l \in L$, and a cell C_k . If $\exists \tau \in [2^l]$ such that $\phi(v) + \tau \cdot 1_d \in C_k$,

- ① $\tau_1 = \min_{\tau \in [2^l]: \phi(v) + \tau \cdot 1_d \in C_k} \tau$
- ② $\tau_2 = \max_{\tau \in [2^l]: \phi(v) + \tau \cdot 1_d \in C_k} \tau$
- ③ $a \leftarrow (k-1)2^l + \sum_{j=0}^{l-1} 2^j |C_j|$
- ④ $I_v \leftarrow I_v \cup \{([a + \tau_1, a + \tau_2], d)\}$

Implicit Preconditioner

Let $x = \mathbf{W}f$. Find the optimal x^* such that

$$\begin{aligned} x^* &= \min_{x \in \mathbb{R}^m} \|x\|_1 \\ \text{s.t. } \mathbf{A}\mathbf{W}^{-1}x &= b. \end{aligned}$$

Implicit Preconditioner

Let $x = \mathbf{W}f$. Find the optimal x^* such that

$$\begin{aligned} x^* &= \min_{x \in \mathbb{R}^m} \|x\|_1 \\ \text{s.t. } \mathbf{A}\mathbf{W}^{-1}x &= b. \end{aligned}$$

Theorem

Given an undirected graph $G = (V, E, w)$ and a mapping $\phi(v) : V \rightarrow [\Delta]^d$ such that

$$\forall u, v \in V, d_G(u, v) \leq \|\phi(u) - \phi(v)\|_1 \leq \alpha \cdot d_G(u, v),$$

we can *efficiently* compute a compressed representation $I = (I_1, I_2, \dots, I_n)$ of a matrix P with full column rank and

- 1 $\kappa(P\mathbf{A}\mathbf{W}^{-1}) \leq \mathcal{O}(\alpha Ld)$
- 2 each I_i of size at most $(d+1)L$

Fast Operations

How can we do matrix-vector products with our compressed representation?

$MatVec(I = (I_1, I_2, \dots, I_n), g \in \mathbb{R}^n)$

① $S \leftarrow \emptyset, \hat{I} \leftarrow \emptyset$

② For $i \in [n] : g_i \neq 0$

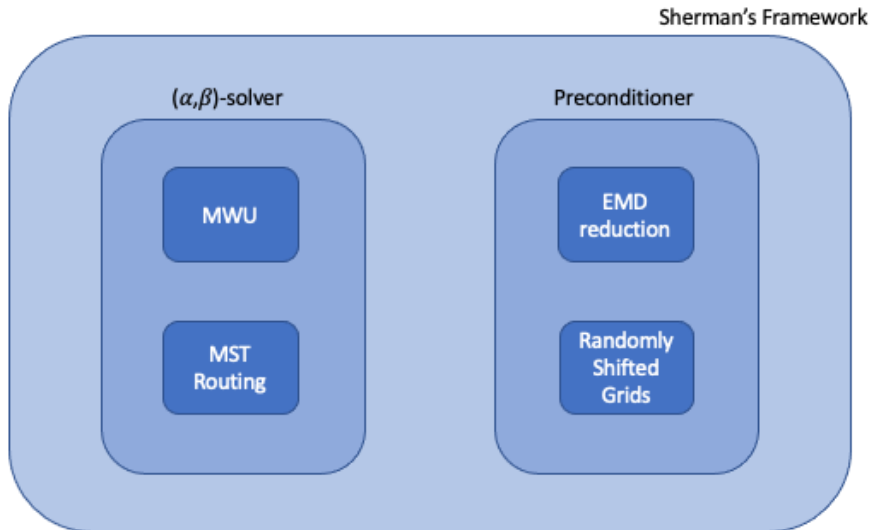
① for each $([a, b], c) \in I_i, S \leftarrow S \cup \{(a, cg_i), (b+1, -cg_i)\}$

③ Sort $S = \{(q_1, z_1), (q_2, z_2), \dots, (q_k, z_k)\}$ such that
 $q_1 \leq q_2 \leq \dots \leq q_k$

④ For each $j \in \{2, 3, \dots, k\} : q_j > q_{j-1}, \hat{I} \leftarrow$
 $\hat{I} \cup \{([q_{j-1}, q_j - 1], \sum_{t: q_t < q_j} z_t)\}$

Sherman's Framework

We use Sherman's framework to solve uncapacitated min-cost flow.



Summary

- 1 Low Hop Emulators and Applications
- 2 Constructing a Low Hop Emulator
- 3 Uncapacitated Min-Cost Flow in Sherman's Framework
- 4 Constructing a Good Preconditioner

Thanks. Questions?